

---

# CS150A Homework 1 – Writing

---

School of Information Science and Technology

April 17, 2024

## 1 SQL (40 PTS)

The following relations keep track of airline flight information:

Flights (fno: integer, distance: integer, price: integer, aid: integer)

Aircraft (aid: integer, aname: string, cruisingrange: integer)

Certified (eid: integer, aid: integer)

Employees (eid: integer, ename: string, salary: integer)

Note that the underlined attributes are the primary keys; the Employees relation describes pilots and other kinds of employees as well; every pilot is certified for some aircraft, and only pilots are certified to fly; each aircraft may operate multiple flights. Answer the following questions.

1. Write a SQL to find the eid, name and salary of employees who are not pilots and sort them from highest to lowest salary.

```
SELECT E.eid, E.name, E.salary
FROM Employees E
WHERE NOT EXISTS
(SELECT *
FROM Certified C
WHERE C.eid = E.eid)
ORDER BY E.salary DESC;
```

2. Write a SQL to find the employee ID (i.e. eid), name (i.e. ename) and salary of all pilots whose salary is greater than the average salary of all pilots, and sort the query results from low to high salary after removing duplicates.

```

WITH Pilots(pid, pname, salary) AS
(SELECT DISTINCT E.eid, E.ename, E.salary
FROM Employees E, Certified C
WHERE E.eid = C.eid)
SELECT pid, pname, salary
FROM Pilots
WHERE salary >
(SELECT AVG(salary)
FROM Pilots)
ORDER BY salary ASC;

```

3. Which of the following queries returns the flight number (i.e. flno) of the flights with the highest price per kilometer (i.e. price / distance)? (The correct answer may be more than one)
- A. `SELECT F.flno`  
`FROM Flights F`  
`WHERE (F.price/F.distance) = MAX(F.price/F.distance);`
  - B. `SELECT F.flno`  
`FROM Flights F`  
`WHERE (F.price/F.distance) >= ALL (SELECT F.price/F.distance FROM Flights F) ;`
  - C. `SELECT F.flno`  
`FROM Flights F`  
`WHERE (F.price/F.distance) = (SELECT MAX(F.price/F.distance) FROM Flights F);`
  - D. `SELECT F.flno, MAX(F.price/F.distance)`  
`FROM Flights F`  
`GROUP BY F.flno`
  - E. None of the above

## 2 INDEX AND B+ TREES (30 PTS)

1. Consider the B+ tree index of order  $d = 2$  shown below:

1. Show the tree that would result from inserting a data entry with key 3 into this tree.
2. Show the B+ tree that would result from deleting the data entry with key 8 from the original tree, assuming that the right sibling is checked for possible redistribution.




Figure 2.1: B+ Tree

2. Suppose that all nodes in our B+ tree have an order of 1500. What's the MAXIMUM number of records we can index with a B+ tree of height 2? (Assume our B+ trees are laid out as in lecture.)

$2.7 * 10^{10}$  records

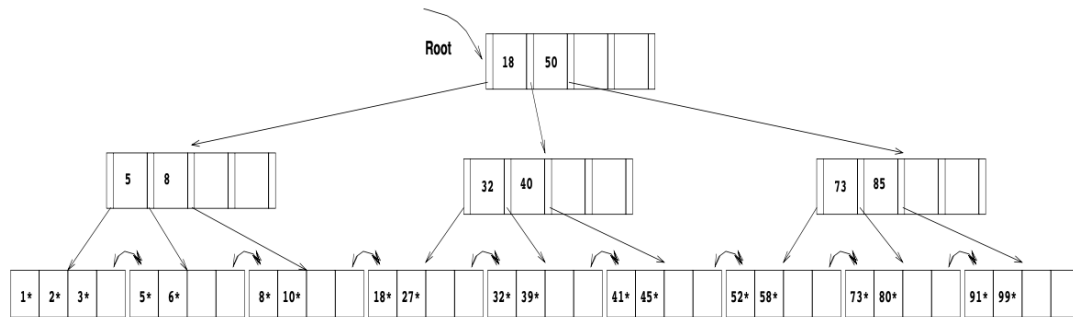


Figure 2.2: answer for 2.1




Figure 2.3: answer for 2.2

### 3 FILE ORGANIZATION (30 PTS)

Consider the following relations:

Emp(eid : integer, *ename* : varchar, *sal* : integer, *age* : integer, *did* : integer)

Dept(did : integer, *budget* : integer, *floor* : integer, *mgr\_eid* : integer)

Salaries range from 10,000 to 100,000, ages vary from 20 to 80, each department has about five employees on average, there are 10 floors, and budgets vary from \$10,000 to \$1 million. You can assume uniform distributions of values.

For each of the following queries, which of the listed index choices would you choose to speed up the query? If your database system does not consider index-only plans (i.e., data records are always retrieved even if enough information is available in the index entry), how would your answer change? Explain briefly.

1. Query: *Print ename, age, and sal for all employees.*
  - A. Clustered hash index on  $\langle ename, age, sal \rangle$  fields of Emp.
  - B. Unclustered hash index on  $\langle ename, age, sal \rangle$  fields of Emp.
  - C. Clustered B+ tree index on  $\langle ename, age, sal \rangle$  fields of Emp.
  - D. Unclustered hash index on  $\langle eid, did \rangle$  fields of Emp.
  - E. No index.
2. Query: *Find the dids of departments that are on the 10th floor and have a budget of less than \$15,000.*
  - A. Clustered hash index on the *floor* field of Dept.
  - B. Unclustered hash index on the *floor* field of Dept.
  - C. Clustered B+ tree index on  $\langle floor, budget \rangle$  fields of Dept.
  - D. Clustered B+ tree index on the *budget* field of Dept.
  - E. No index.

Assume that each page in our system can hold 128 KB (1 KB = 1024bytes), integers are 32-bits wide, and bytes are 8-bits wide.

Consider the following relation:

```
CREATE TABLE Submissions(  
  record_id integer UNIQUE,  
  assignment_id integer,  
  student_id integer,  
  time_submitted integer,  
  grade_received byte,  
  PRIMARY KEY (assignment_id, student_id)  
);
```

Assume the column record\_id corresponds to the row's actual record ID.

1. How large (in bytes) is a record?

We simply add up the sizes of each field in a record. We have 4 integer fields and 1 byte field, which is  $4 \times 4 + 1 = 17$  bytes.

2. Suppose we begin each page with a 32-bytes header plus a bitmap. At most, how many records can fit in an unpacked page?

Let's convert everything into bits. First, a page holds  $1,024 \times 64 \times 8 = 1,048,576$  bits while a record holds  $17 \times 8 = 136$  bits. Now, remember that in an unpacked page, each record needs an additional bit to represent whether or not it is valid (e.g. has been deleted). This means we need 1 more bit per record, so each record in fact requires 137 bits. Finally, we have an extra 32 bytes ( $32 \times 8 = 256$  bits) reserved for the page header. This gives us a maximum of  $(1,048,576 - 256) / 137 \approx 7,651$  records.

Answer for 3.1:

1. We should create an unclustered hash index on  $\langle \text{ename}, \text{age}, \text{sal} \rangle$  fields of Emp (b) since then we could do an index only scan. If our system does not include index only plans then we shouldn't create an index for this query (e). Since this query requires us to access all the Emp records, an index won't help us any, and so should we access the records using a filescan.
2. We should create a clustered dense B+ tree index (c) on  $\langle \text{floor}, \text{budget} \rangle$  fields of Dept, since the records would be ordered on these fields then. So when executing this query, the first record with floor = 10 must be retrieved, and then the other records with floor = 10 can be read in order of budget. Note that this plan, which is the best for this query, is not an index-only plan (must look up dids).