

# CS150A Database

Lu Sun

School of Information Science and Technology

ShanghaiTech University

Feb. 28, 2024

Today:

- Introduction to SQL
- DDL & DML

Readings:

- Database Management Systems (DBMS), Chapter 5
- Lecture note SQL I

*Acknowledgement: Joe Hellerstein@UCBerkeley's scourse notes*

# SQL Roots

- Developed @IBM Research in the 1970s
  - System R project
  - Vs. Berkeley's Quel language (Ingres project)
- Commercialized/Popularized in the 1980s
  - "Intergalactic Dataspeak"
  - IBM beaten to market by a startup called Oracle

# SQL's Persistence

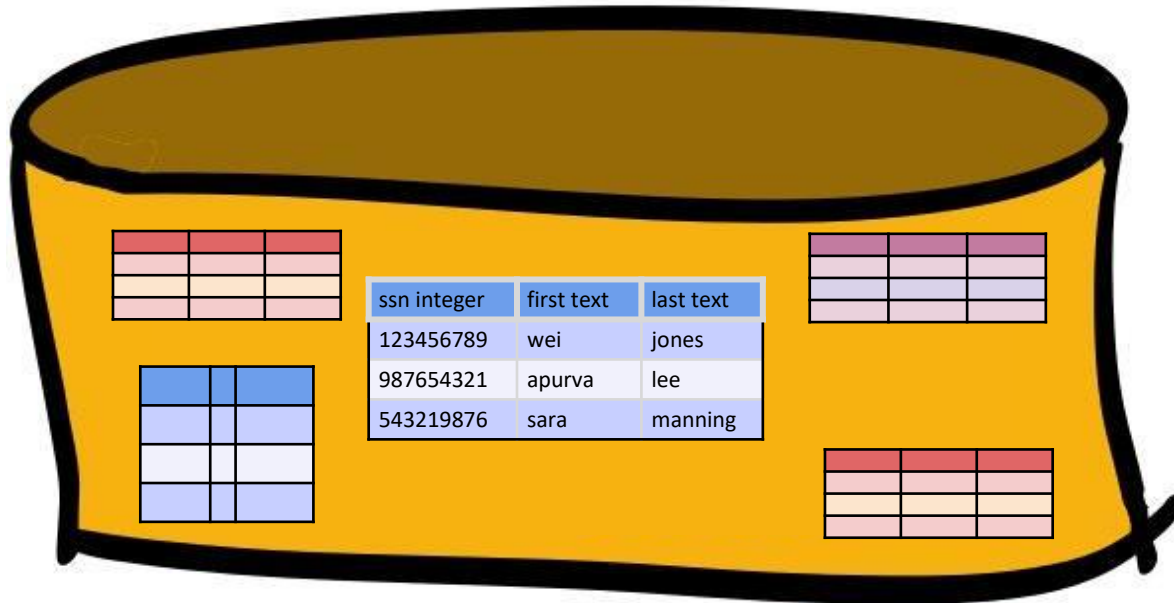
- Over 40 years old!
- Questioned repeatedly
  - 90's: Object-Oriented DBMS (OQL, etc.)
  - 2000's: XML (Xquery, Xpath, XSLT)
  - 2010's: NoSQL & MapReduce
- SQL keeps re-emerging as the standard
  - Even Hadoop, Spark etc. mostly used via SQL
  - May not be perfect, but it is useful

# SQL Pros and Cons

- Declarative!
  - Say *what* you want, not *how* to get it
- Implemented widely
  - With varying levels of efficiency, completeness
- Constrained
  - Not targeted at Turing-complete tasks
- General-purpose and feature-rich
  - many years of added features
  - extensible: callouts to other languages, data sources

# Relational Terminology

- *Database*: Set of named Relations



# Relational Terminology, Pt 2.

- *Database*: Set of named Relations
- *Relation (Table)*:
  - *Schema*: description ( “metadata” )
  - *Instance*: set of data satisfying the schema

ssn integer	first text	last text
123456789	wei	jones
987654321	apurva	lee
543219876	sara	manning

# Relational Terminology, Pt. 3

- *Database*: Set of named Relations
- *Relation (Table)*:
  - *Schema*: description ( “metadata” )
  - *Instance*: set of data satisfying the schema
- *Attribute (Column, Field)*

first text
wei
apurva
sara

# Relational Terminology, Pt. 4

- *Database*: Set of named Relations
- *Relation* (Table):
  - *Schema*: description ( “metadata” )
  - *Instance*: set of data satisfying the schema
- *Attribute* (Column, Field)
- *Tuple* (Record, Row)

543219876	sara	manning
-----------	------	---------



# Relational Tables

- *Schema* is fixed:
  - unique attribute names, *atomic* types
  - folks (ssn integer, first text, last text)
- *Instance* can change often
  - a *multiset* of “rows” ( “tuples” )

```
{(123456789, ' wei' , ' jones' ),  
(987654321, ' apurva' , ' lee' ),  
(543219876, 'sara' , 'manning' ),  
(987654321, ' apurva' , ' lee' )}
```

# Quick Check 1

- Is this a relation?

num integer	street text	zip integer
84	Maple Ave	54704
22	High	Street
75	Hearst Ave	94720

76425

# Quick Check 2

- Is this a relation?

num integer	street text	num integer
84	Maple Ave	54704
22	High Street	76425
75	Hearst Ave	94720

# Quick Check 3

- Is this a relation?

first text	last text	addr address
wei	jones	(84, 'Maple', 54704)
apurva	lee	(22, 'High', 76425)
sara	manning	(75, 'Hearst', 94720)

# SQL Language

- Two sublanguages:
  - DDL – Data Definition Language
    - Define and modify schema
  - DML – Data Manipulation Language
    - Queries can be written intuitively.
- RDBMS responsible for efficient evaluation.
  - Choose and run algorithms for declarative queries
    - Choice of algorithm must not affect query answer.

# Example Database

## Sailors

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

## Boats

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

## Reserves

<u>sid</u>	bid	day
1	102	9/12/2015
2	102	9/13/2015

# The SQL DDL: Sailors

```
CREATE TABLE Sailors (  
    sid INTEGER,  
    sname CHAR(20),  
    rating INTEGER,  
    age FLOAT
```

<b><u>sid</u></b>	<b>sname</b>	<b>rating</b>	<b>age</b>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

# The SQL DDL: Sailors, Pt. 2

```
CREATE TABLE Sailors (  
    sid INTEGER,  
    sname CHAR(20),  
    rating INTEGER,  
    age FLOAT  
    PRIMARY KEY (sid));
```

<b><u>sid</u></b>	<b>sname</b>	<b>rating</b>	<b>age</b>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27



# The SQL DDL: Primary Keys

```
CREATE TABLE Sailors (  
    sid INTEGER,  
    sname CHAR(20),  
    rating INTEGER,  
    age FLOAT
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

- Primary Key column(s)
  - Provides a unique “lookup key” for the relation
  - Cannot have any duplicate values
  - Can be made up of >1 column
    - E.g. (firstname, lastname)

# The SQL DDL: Boats

```
CREATE TABLE Sailors (  
    sid INTEGER,  
    sname CHAR(20),  
    rating INTEGER,  
    age FLOAT
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

```
CREATE TABLE Boats (  
    bid INTEGER,  
    bname CHAR (20),  
    color CHAR(10),  
    PRIMARY KEY (bid));
```

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

# The SQL DDL: Reserves

```
CREATE TABLE Sailors (  
  sid INTEGER,  
  sname CHAR(20),  
  rating INTEGER,  
  age FLOAT
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

```
CREATE TABLE Boats (  
  bid INTEGER,  
  bname CHAR(20),  
  color CHAR(10),  
  PRIMARY KEY (bid));
```

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

```
CREATE TABLE Reserves (  
  sid INTEGER,  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (sid, bid, day);
```

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

# The SQL DDL: Reserves Pt. 2

```
CREATE TABLE Sailors (  
  sid INTEGER,  
  sname CHAR(20),  
  rating INTEGER,  
  age FLOAT
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

```
CREATE TABLE Boats (  
  bid INTEGER,  
  bname CHAR(20),  
  color CHAR(10),  
  PRIMARY KEY (bid));
```

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

```
CREATE TABLE Reserves (  
  sid INTEGER,  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (sid, bid, day),  
  FOREIGN KEY (sid) REFERENCES Sailors,
```

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

# The SQL DDL: Foreign Keys

```
CREATE TABLE Sailors (  
  sid INTEGER,  
  sname CHAR(20),  
  rating INTEGER,  
  age FLOAT
```

```
CREATE TABLE Boats (  
  bid INTEGER,  
  bname CHAR(20),  
  color CHAR(10),  
  PRIMARY KEY (bid));
```

```
CREATE TABLE Reserves (  
  sid INTEGER,  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (sid, bid, day),  
  FOREIGN KEY (sid) REFERENCES Sailors,  
  FOREIGN KEY (bid) REFERENCES Boats);
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

# The SQL DDL: Foreign Keys Pt. 2

- Foreign key references a table
  - Via the primary key of that table
- Need not share the name of the referenced primary key

```
CREATE TABLE Reserves (  
  sid INTEGER,  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (sid, bid, day),  
  FOREIGN KEY (sid) REFERENCES  
    Sailors,  
  FOREIGN KEY (bid) REFERENCES  
    Boats);
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

# The SQL DML

- Find all 27-year-old sailors:  

```
SELECT *  
FROM Sailors AS S  
WHERE S.age=27;
```

- To find just names and rating,  
replace the first line to:  

```
SELECT S.sname,  
       S.rating
```

## Sailors

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

# Basic Single-Table Queries

- `SELECT` [*DISTINCT*] *<column expression list>*  
`FROM` *<single table>*  
[`WHERE` *<predicate>*]
- Simplest version is straightforward
  - Produce all tuples in the table that satisfy the predicate
  - Output the expressions in the `SELECT` list
  - Expression can be a column reference, or an arithmetic expression over column refs



# SELECT DISTINCT

```
SELECT DISTINCT S.name, S.gpa  
FROM students S  
WHERE S.dept = 'CS'
```

- DISTINCT specifies removal of duplicate rows before output
- Can refer to the students table as “S”, this is called an alias

# ORDER BY

- ```
SELECT S.name, S.gpa, S.age*2 AS a2
FROM Students S
WHERE S.dept = 'CS'
ORDER BY S.gpa, S.name, a2;
```
- ORDER BY clause specifies output to be sorted
  - *Lexicographic* ordering
- Obviously must refer to columns in the output
  - Note the AS clause for naming output columns!

# ORDER BY, Pt. 2

- ```
SELECT  S.name, S.gpa, S.age*2 AS a2
FROM    Students S
WHERE   S.dept = 'CS'
ORDER BY S.gpa DESC, S.name ASC, a2;
```
- Ascending order by default, but can be overridden
  - DESC flag for descending, ASC for ascending
  - Can mix and match, lexicographically

# LIMIT

- ```
SELECT  S.name, S.gpa, S.age*2 AS a2
FROM    Students S
WHERE   S.dept = 'CS'
ORDER BY S.gpa DESC, S.name ASC, a2;
LIMIT 3 ;
```
- Only produces the first <integer> output rows
- Typically used with ORDER BY
  - Otherwise the output is *non-deterministic*
  - Not a “pure” declarative construct in that case – output set depends on algorithm for query processing

# Aggregates

- `SELECT [DISTINCT] AVG(S. gpa)`  
`FROM Students S`  
`WHERE S. dept = 'CS'`
- Before producing output, compute a summary (a.k.a. an *aggregate*) of some arithmetic expression
- Produces 1 row of output
  - with one column in this case
- Other aggregates: SUM, COUNT, MAX, MIN

# GROUP BY

```
SELECT [DISTINCT] AVG(S. gpa), S. dept  
FROM Students S  
GROUP BY S. dept
```

- Partition table into groups with same GROUP BY column values
  - Can group by a list of columns
- Produce an aggregate result per group
  - Cardinality of output = # of distinct group values
- Note: can put grouping columns in SELECT list

# HAVING

```
SELECT [DISTINCT] AVG(S. gpa), S. dept  
FROM Students S  
GROUP BY S. dept  
HAVING COUNT(*) > 2
```

- The HAVING predicate filters groups
- HAVING is applied *after* grouping and aggregation
  - Hence can contain anything that could go in the SELECT list
  - I.e. aggs or GROUP BY columns
- HAVING can only be used in aggregate queries
- It' s an optional clause

# Putting it all together

```
SELECT S.dept, AVG(S.gpa), COUNT(*)  
FROM Students S  
WHERE S.gender = 'F'  
GROUP BY S.dept  
HAVING COUNT(*) >= 2  
ORDER BY S.dept;
```



# DISTINCT Aggregates

Are these the same or different?

```
SELECT COUNT(DISTINCT S.name)
FROM Students S
WHERE S.dept = 'CS' ;
```

```
SELECT DISTINCT COUNT(S.name)
FROM Students S
WHERE S.dept = 'CS' ;
```

# What Is This Asking For?

```
SELECT S.name, AVG(S.gpa)
FROM Students S
GROUP BY S.dept;
```

# SQL DML:

## General Single-Table Queries

- **SELECT [DISTINCT]** *<column expression list>*  
**FROM** *<single table>*  
**[WHERE** *<predicate>*  
**[GROUP BY** *<column list>*  
**[HAVING** *<predicate>*] ]  
**[ORDER BY** *<column list>*]  
**[LIMIT** *<integer>*];

# Summary

- Relational model has **well-defined query semantics**
- Modern SQL extends “pure” relational model  
*(some extra goodies for duplicate row, non-atomic types...  
more in next lecture)*
- Typically, many ways to write a query
  - DBMS figures out a fast way to execute a query, regardless of how it is written.