

CS150A Database

Lu Sun

School of Information Science and Technology

ShanghaiTech University

Apr. 10, 2024

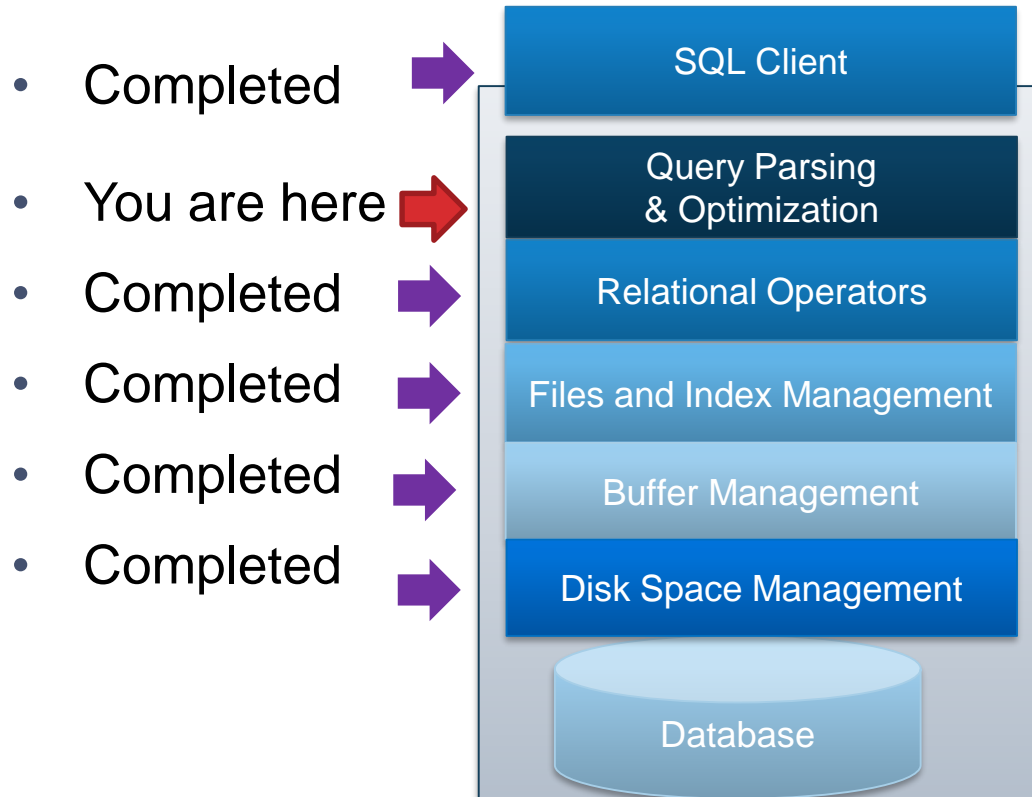
Today:

- Query Optimization I:
 - The Plan Space

Readings:

- Database Management Systems (DBMS), Chapter 15

Architecture of a DBMS



Query Optimization is Magic

- The bridge between a *declarative* domain-specific language...
 - “What” you want as an answer
- ... and custom *imperative* computer programs
 - “How” to compute the answer
- In 2018 terms:
 - This is AI-driven Software Synthesis
 - That’s not just marketing!
 - Similar to cutting edge AI work today
 - Optimization + heuristic pruning
 - Research exploring the use of modern AI techniques to improve that pruning (e.g. Deep Reinforcement Learning)

Invented in 1979 by Pat Selinger et al.

- We'll focus on “System R” (“Selinger”) optimizers
- “Cascades” optimizer is the other common one
 - Later, with notable differences, but similar big picture

Access Path Selection
in a Relational Database Management System

P. Griffiths Selinger
M. M. Astrahan
D. D. Chamberlin
R. A. Lorie
T. G. Price

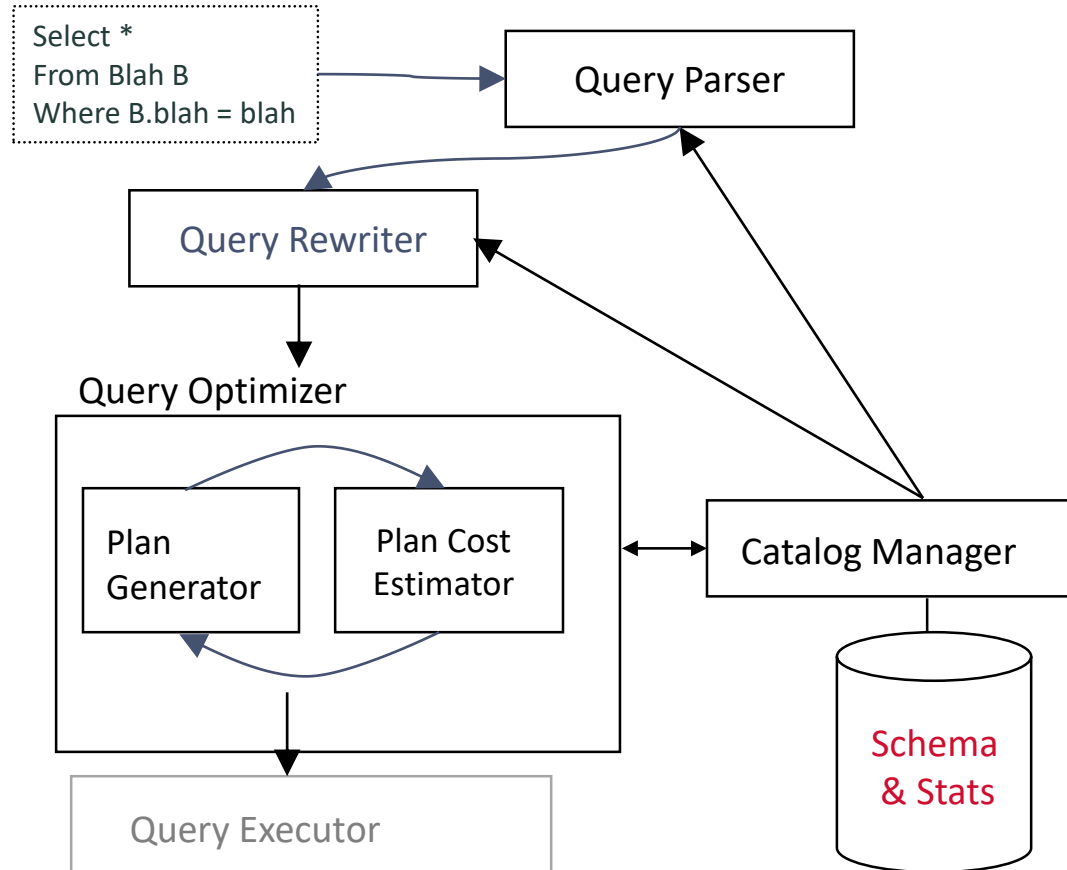
IBM Research Division, San Jose, California 95193



ABSTRACT: In a high level query and data manipulation language such as SQL, requests are submitted sequentially, without

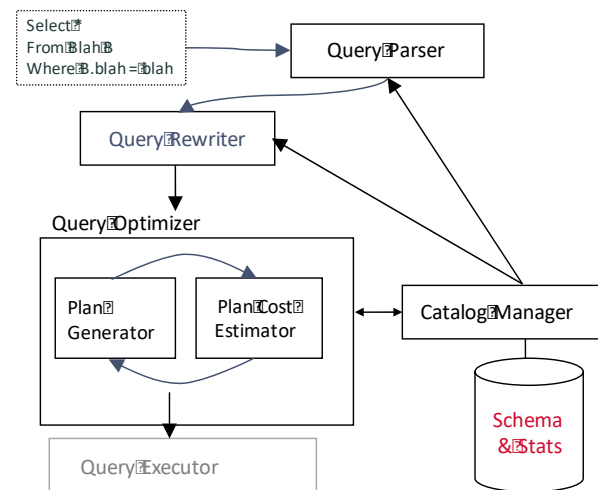
retrieval. Nor does a user specify in what order joins are to be performed. The System R optimizer chooses both join order

Query Parsing & Optimization



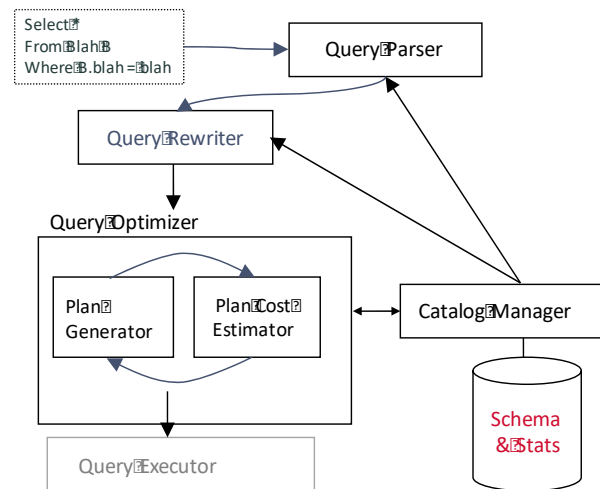
Query Parsing & Optimization Part 2

- Query parser
 - Checks correctness, authorization
 - Generates a parse tree
 - Straightforward



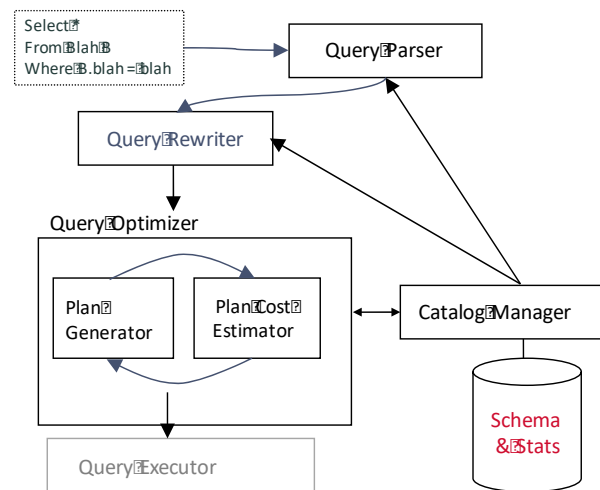
Query Parsing & Optimization Part 3

- Query rewriter
 - Converts queries to canonical form
 - flatten views
 - subqueries into fewer query blocks
 - Weak spot in many open-source DBMSs



Query Parsing & Optimization Part 4

- “Cost-based” Query Optimizer
 - Optimizes 1 query block at a time
 - Select, Project, Join
 - GroupBy/Agg
 - Order By (if top-most block)
 - Uses catalog stats to find least-“cost” plan per query block
 - “Soft underbelly” of every DBMS
 - Sometimes not truly “optimal”

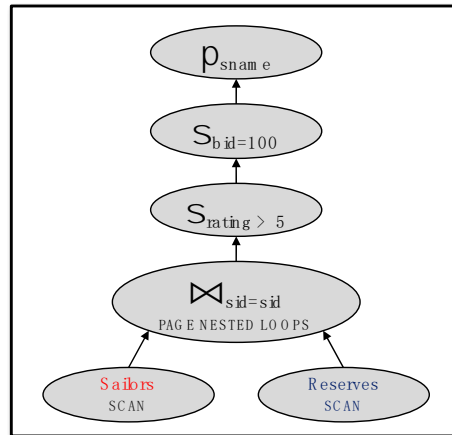


Query Optimization Overview

- Query block can be converted to relational algebra
- Rel. Algebra converts to tree
- Each operator has implementation choices
- Operators can also be applied in different orders!

```
SELECT S.sname  
  FROM Reserves R, Sailors S  
 WHERE R.sid=S.sid  
       AND R.bid=100  
       AND S.rating>5
```

$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)}$
 $(Reserves \bowtie Sailors)$



Query Optimization: The Components

- Three beautifully orthogonal concerns:
 - Plan space:
 - for a given query, what plans are considered?
 - Cost estimation:
 - how is the cost of a plan estimated?
 - Search strategy:
 - how do we “search” in the “plan space”?

Query Optimization: The Goal

- Optimization goal:
 - Ideally: Find the plan with least actual cost.
 - Reality: Find the plan with least estimated cost.
 - And try to avoid really bad actual plans!

Today

- We will get a feel for the plan space
- Explore one simple example query

Relational Algebra Equivalences: Selections

- Selections:
 - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots(\sigma_{cn}(R))\dots)$ (cascade)
 - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$ (commute)

Relational Algebra Equivalences: Projections

- Selections:
 - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots(\sigma_{cn}(R))\dots)$ (cascade)
 - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$ (commute)
- Projections:
 - $\pi_{a1}(R) \equiv \pi_{a1}(\dots(\pi_{a1, \dots, an-1}(R))\dots)$ (cascade)

Relational Algebra Equivalences: Cartesian Product

- Selections:
 - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots(\sigma_{cn}(R))\dots)$ (cascade)
 - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$ (commute)
- Projections:
 - $\pi_{a1}(R) \equiv \pi_{a1}(\dots(\pi_{a1, \dots, an-1}(R))\dots)$ (cascade)
- Cartesian Product
 - $R \times (S \times T) \equiv (R \times S) \times T$ (associative)
 - $R \times S \equiv S \times R$ (commutative)



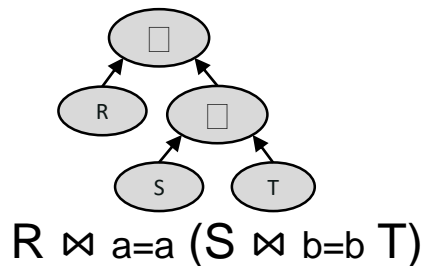
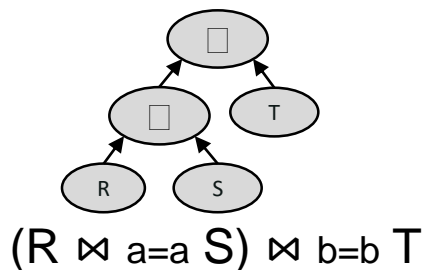
Are Joins Associative and Commutative?

- After all, just Cartesian Products with Selections
- You can think of them as associative and commutative...
- ...But beware of join turning into cross-product!
 - Consider $R(a,z)$, $S(a,b)$, $T(b,y)$
 - $(S \bowtie_{b=b} T) \bowtie_{a=a} R \not\equiv S \bowtie_{b=b} (T \bowtie_{a=a} R)$ (*not legal!!*)
 - $(S \bowtie_{b=b} T) \bowtie_{a=a} R \not\equiv S \bowtie_{b=b} (T \times R)$ (*not the same!!*)
 - $(S \bowtie_{b=b} T) \bowtie_{a=a} R \alpha S \bowtie_{b=b \wedge a=a} (T \times R)$ (*the same!!*)

Join ordering, again

- Similarly, note that some join orders have cross products, some don't
- Equivalent for the query above:

```
SELECT *
  FROM R, S, T
 WHERE R.a = S.a
    AND S.b = T.b;
```



Some Common Heuristics: Selections



- Selection cascade and pushdown
 - Apply selections as soon as you have the relevant columns
 - Ex:
 - $\pi_{\text{sname}} (\sigma_{(\text{bid}=100 \wedge \text{rating} > 5)} (\text{Reserves} \bowtie_{\text{sid}=\text{sid}} \text{Sailors}))$
 - $\pi_{\text{sname}} (\sigma_{\text{bid}=100} (\text{Reserves}) \bowtie_{\text{sid}=\text{sid}} \sigma_{\text{rating} > 5} (\text{Sailors}))$

Some Common Heuristics: Projections

- Projection cascade and pushdown
 - Keep only the columns you need to evaluate downstream operators
 - Ex:
 - $\pi_{\text{sname}} \sigma_{(\text{bid}=100 \wedge \text{rating} > 5)} (\text{Reserves} \bowtie_{\text{sid}=\text{sid}} \text{Sailors})$
 - $\pi_{\text{sname}} (\pi_{\text{sid}} (\sigma_{\text{bid}=100} (\text{Reserves})) \bowtie_{\text{sid}=\text{sid}} \pi_{\text{sname}, \text{sid}} (\sigma_{\text{rating} > 5} (\text{Sailors})))$

Some Common Heuristics

- Avoid Cartesian products
 - Given a choice, do theta-joins rather than cross-products
 - Consider $R(a,b)$, $S(b,c)$, $T(c,d)$
 - Favor $(R \bowtie S) \bowtie T$ over $(R \times T) \bowtie S$



Physical Equivalences

- Base table access,
with single-table selections and projections
 - Heap scan
 - Index scan (if available on referenced columns)
- Equijoins
 - Block (Chunk) Nested Loop: simple, exploits extra memory
 - Index Nested Loop: often good if 1 rel small and the other indexed properly
 - Sort-Merge Join: good with small memory, equal-size tables
 - Grace Hash Join: even better than sort with 1 small table
 - Or Hybrid if you have it
- Non-Equijoins
 - Block Nested Loop

Schema for Examples

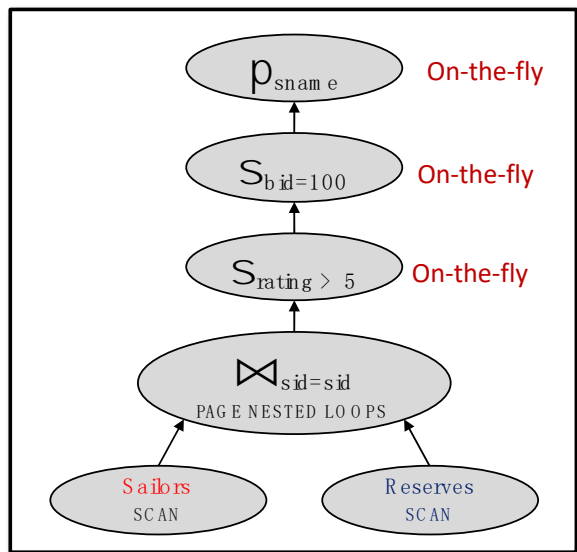
Sailors (sid: integer, sname: text, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: date, rname: text)

- Reserves:
 - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
 - Assume there are 100 boats
- Sailors:
 - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
 - Assume there are 10 different ratings
- Assume we have 5 pages to use for joins.

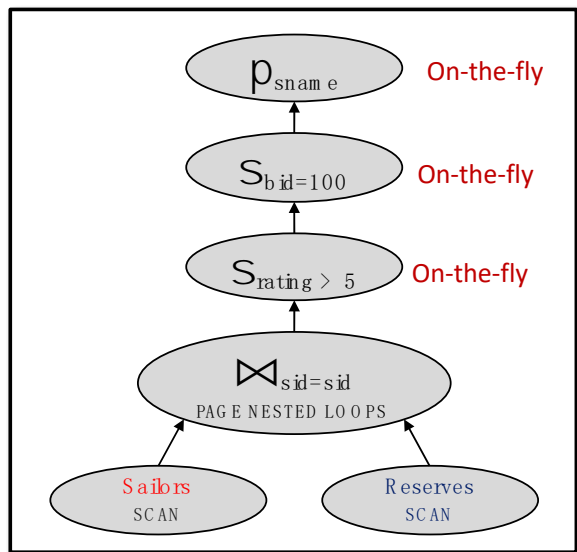
Motivating Example: Plan 1

- Here's a reasonable query plan:



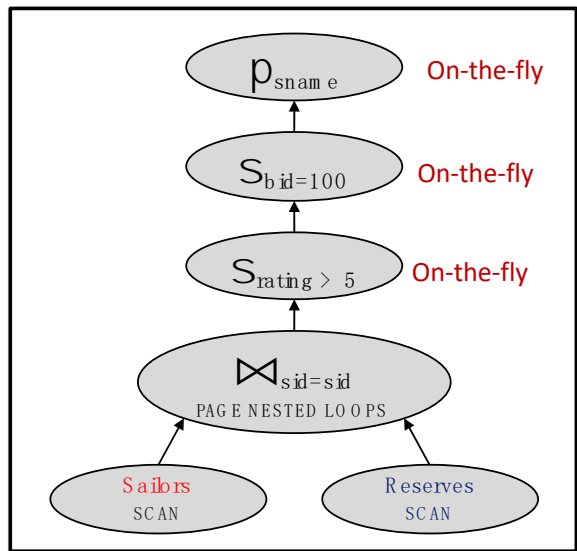
```
SELECT S.sname
FROM Reserves R, sailors S
WHERE R.sid=S.sid
AND R.bid=100
AND S.rating>5
```

Motivating Example: Plan 1 Cost




- Let's estimate the cost:
- Scan Sailors (500 IOs)
- For each page of Sailors, Scan Reserves (1000 IOs)
- Total: $500 + 500 \cdot 1000$
 - **500,500 IOs**

Motivating Example: Plan 1 Cost Analysis

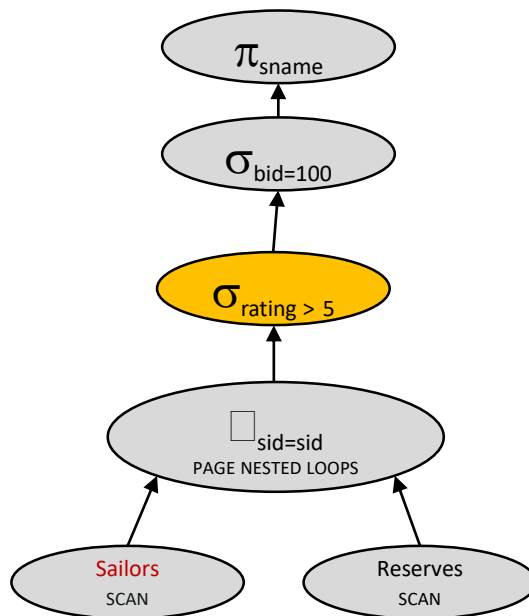


- Cost: $500 + 500 * 1000$ I/Os
- By no means the worst plan!
- Misses several opportunities:
 - selections could be 'pushed' down
 - no use made of indexes
- Goal of optimization:
 - Find faster plans that compute the same answer.

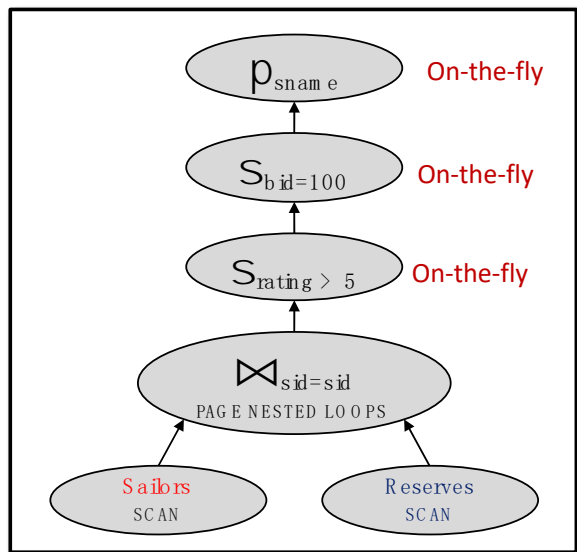
Selection Pushdown



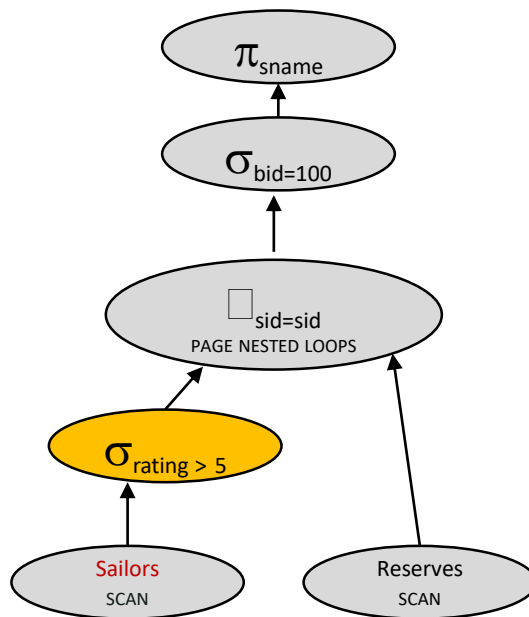
500,500 IOs



Selection Pushdown, cont



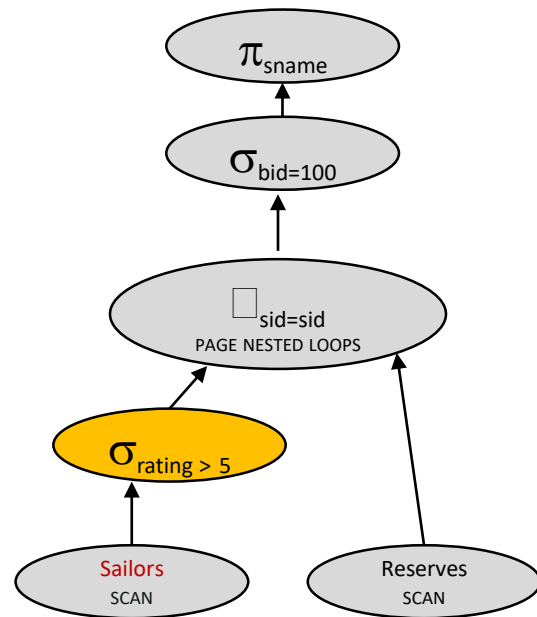
500,500 IOs



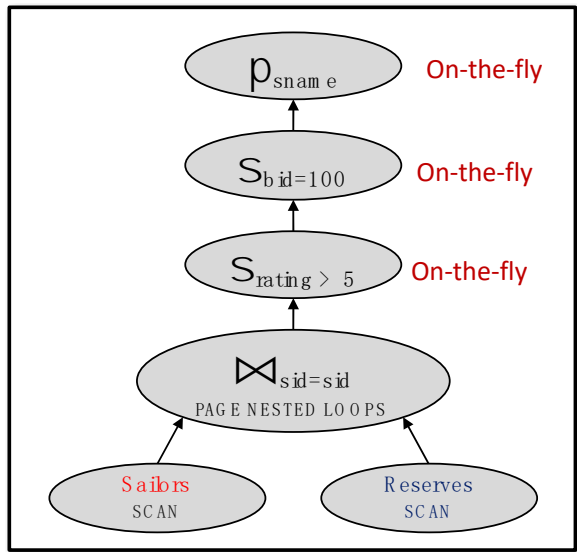
Cost?

Query Plan 2 Cost

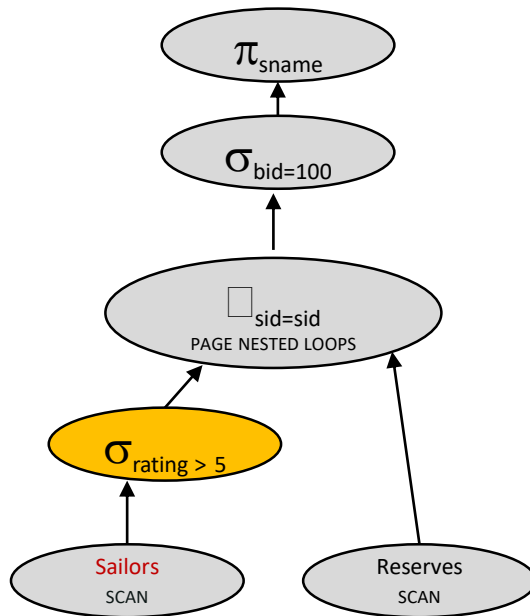
- Let's estimate the cost:
- Scan Sailors (500 IOs)
- For each pageful of high-rated Sailors,
Scan Reserves (1000 IOs)
- Total: $500 + ??? * 1000$
- Total: $500 + 250 * 1000$



Decision?

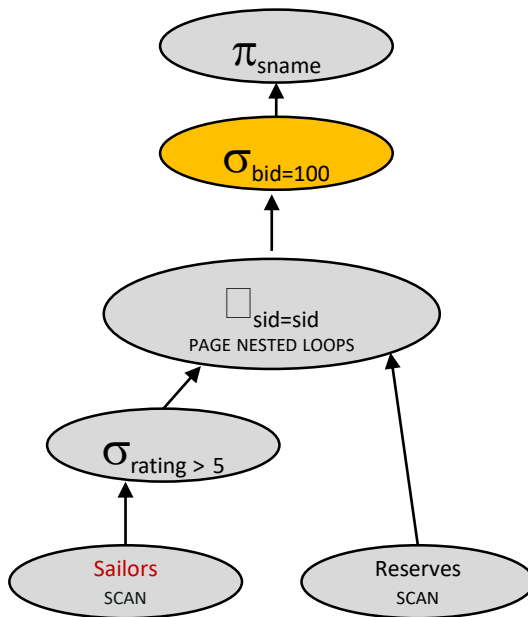
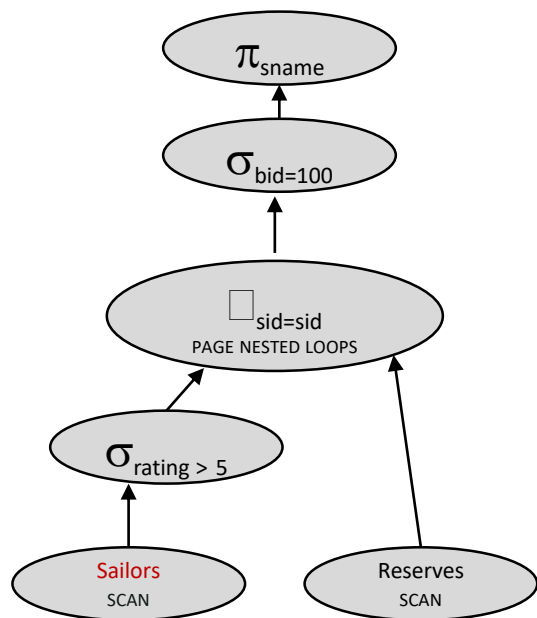


500,500 IOs



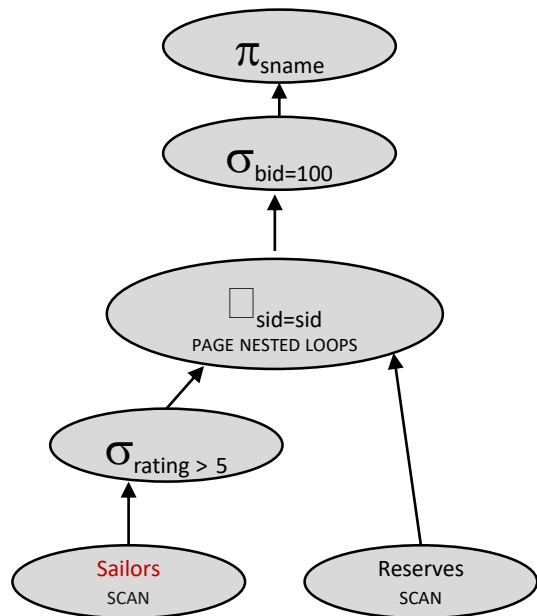
250,500 IOs

More Selection Pushdown

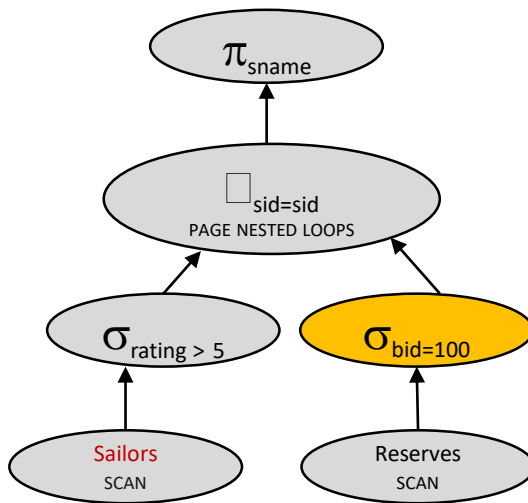


250,500 IOs

More Selection Pushdown, cont



250,500 IOs

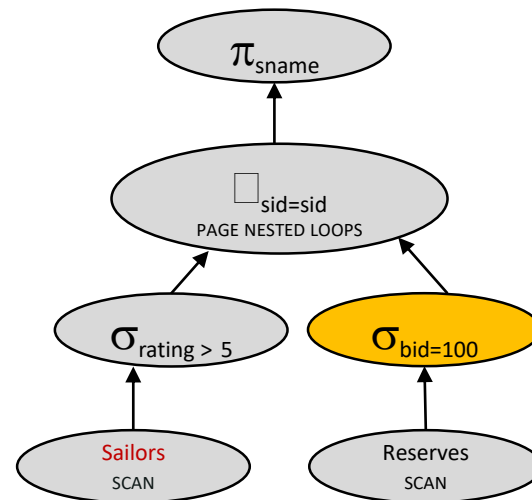


Cost???

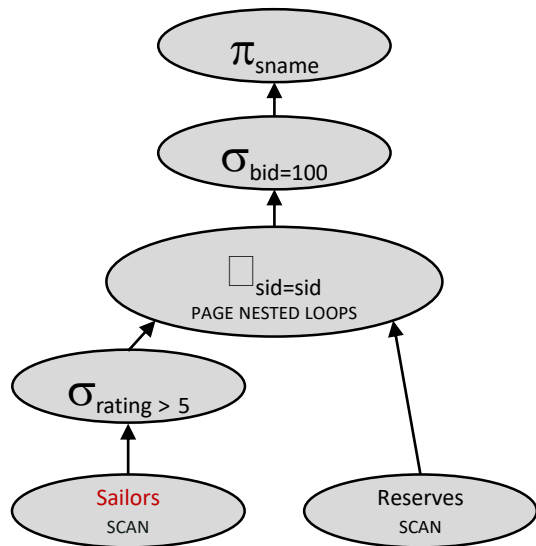
Query Plan 3 Cost Analysis

Let's estimate the cost:


- Scan Sailors (500 IOs)
- For each pageful of high-rated Sailors,
Do what? (??? IOs)
- Total: $500 + 250 * ???$
- Total: $500 + 250 * 1000!$



More Selection Pushdown Analysis



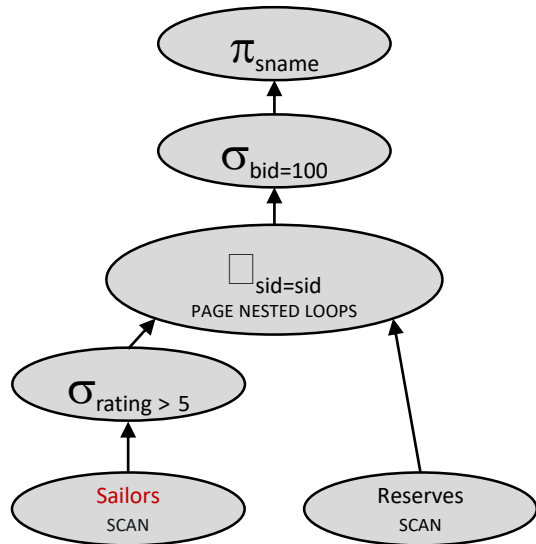
250,500 IOs



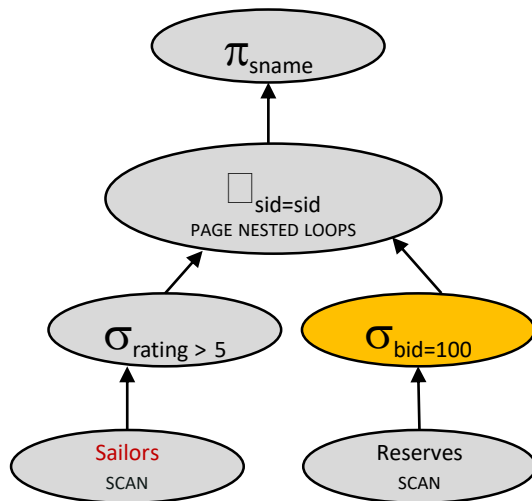
250,500 IOs

Pushing a selection into the inner loop of a nested loop join doesn't save I/Os! Essentially equivalent to having the selection above.

Decision 2





250,500 IOs





250,500 IOs

Join Ordering



250,500 IOs

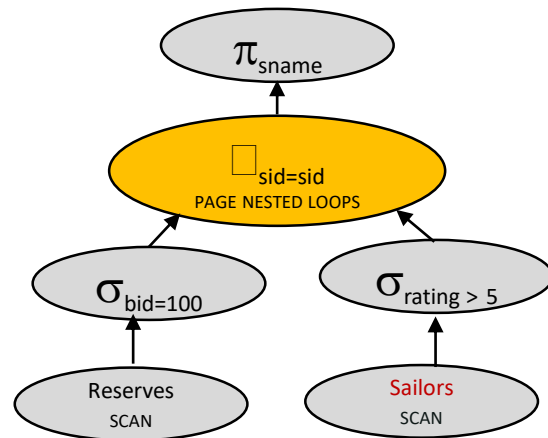
Join Ordering, cont



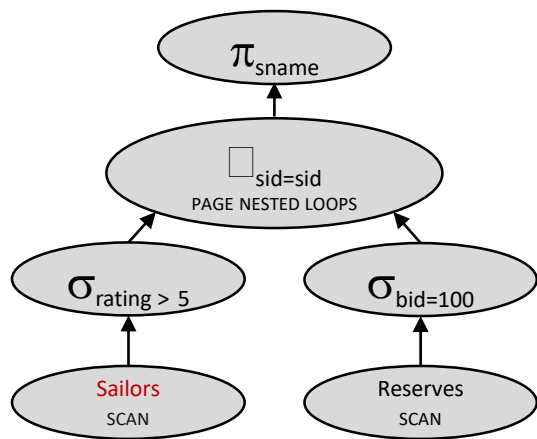
250,500 IOs

Query Plan 4 Cost

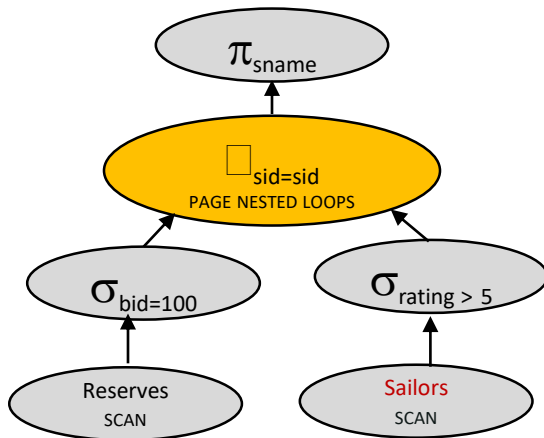
- Let's estimate the cost:
- Scan Reserves (1000 IOs)
- For each pageful of Reserves for bid 100,
Scan Sailors (500 IOs)
- Total: 1000 + ???*500
- Total: **1000 + 10*500**



Decision 3

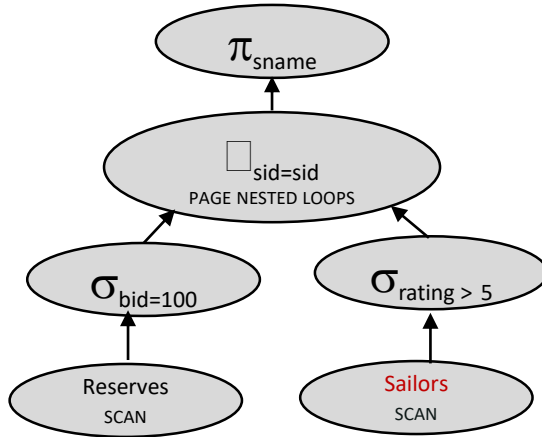


250,500 IOs




6000 IOs

Materializing Inner Loops




6000 IOs




Cost???

Materializing Inner Loops, cont



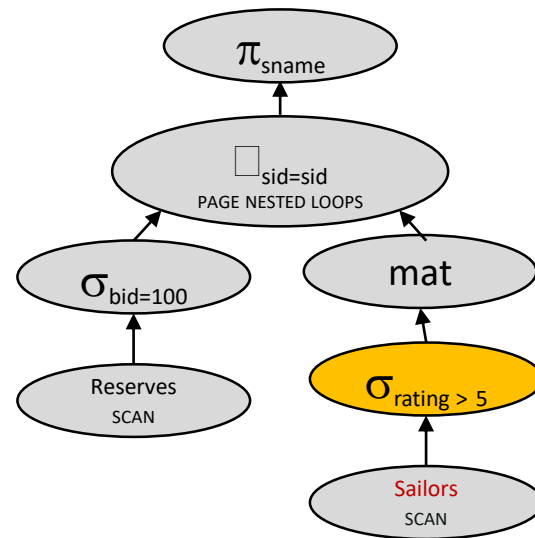
6000 IOs




Cost???

Plan 5 Cost Analysis


- Let's estimate the cost:
- Scan Reserves (1000 IOs)
- Scan Sailors (500 IOs)
- Materialize Temp table T1 (??? IOs)
- For each pageful of Reserves for bid 100,
Scan T1 (??? IOs)
- Total: $1000 + 500 + ??? + 10 * ???$
- $1000 + 500 + 250 + (10 * 250)$



Materializing Inner Loops, cont.




6000 IOs




4250 IOs

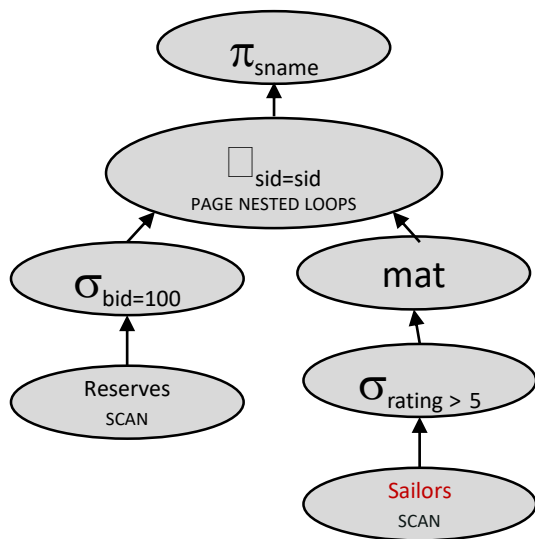
Join Ordering Again



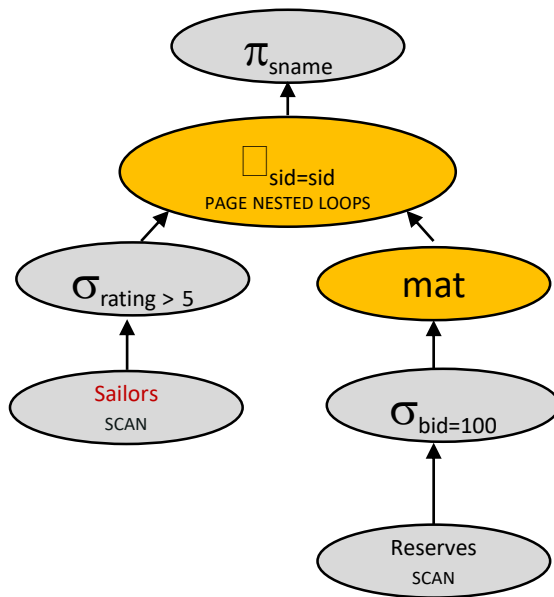
4250 IOs



Join Ordering Again, Cont



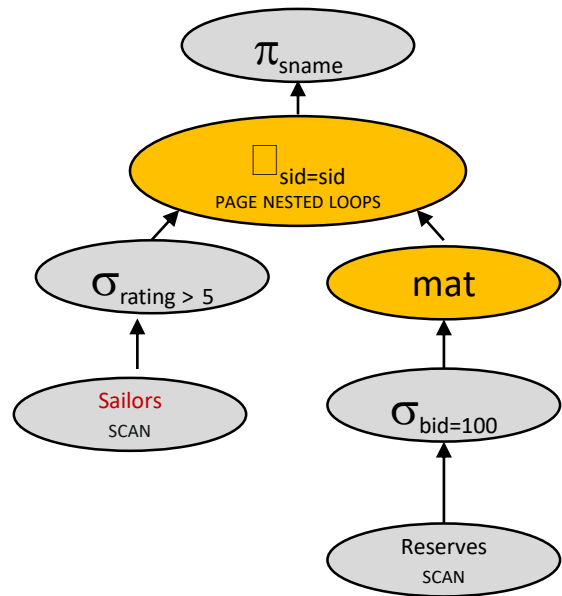
4250 IOs




Cost???

Plan 6 Cost Analysis


- Let's estimate the cost:
- Scan Sailors (500 IOs)
- Scan Reserves (1000 IOs)
- Materialize Temp table T1 (??? IOs)
- For each pageful of high-rated Sailors, Scan T1 (??? IOs)
- Total: $500 + 1000 + ??? + 250 * ???$
- $500 + 1000 + 10 + (250 * 10)$



Decision 4

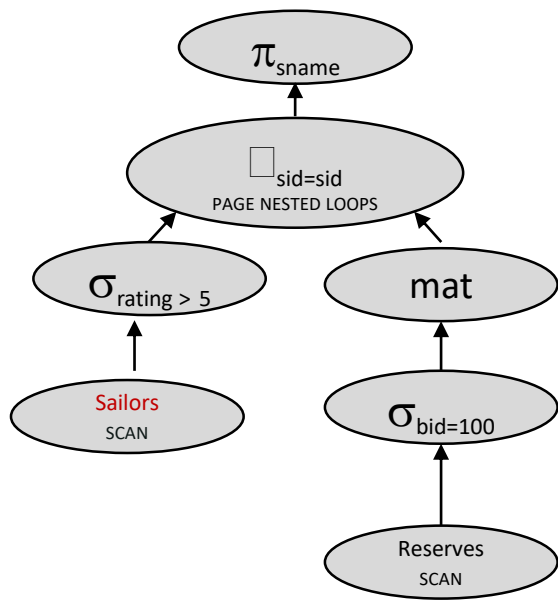


4250 IOs

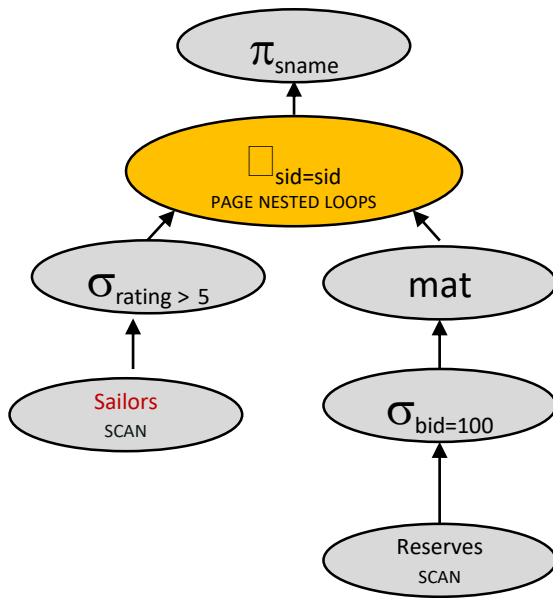


4010 IOs

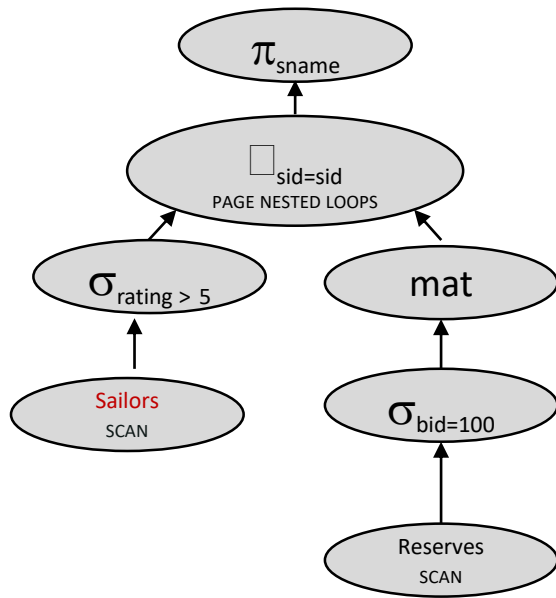
Join Algorithm



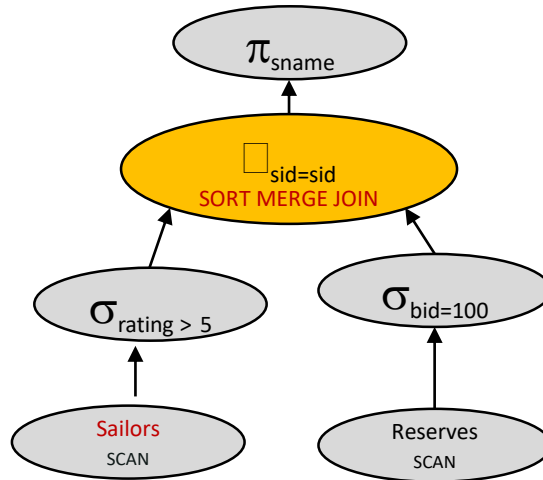
4010 IOs



Join Algorithm, cont.



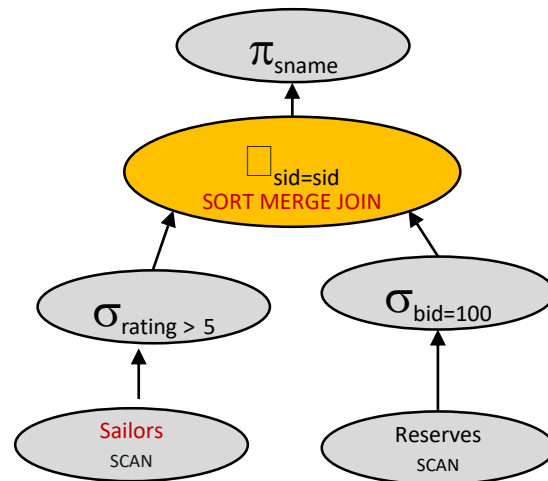
4010 IOs



Cost???

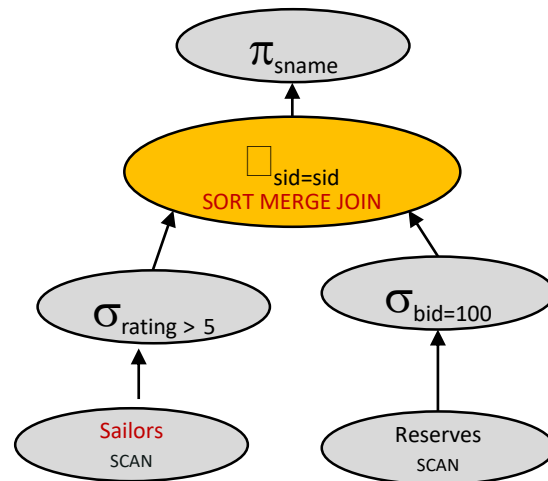
Query Plan 7 Cost Analysis

- With 5 buffers, cost of plan:
- Scan Reserves (1000)
- Scan Sailors (500)
- Sort high-rated sailors (???)
Note: pass 0 doesn't do read I/O, just gets input from select.
- Sort reservations for boat 100 (???)
Note: pass 0 doesn't do read I/O, just gets input from select.
- How many passes for each sort with \log_4 ?
- Merge $(10+250) = 260$
- Total:

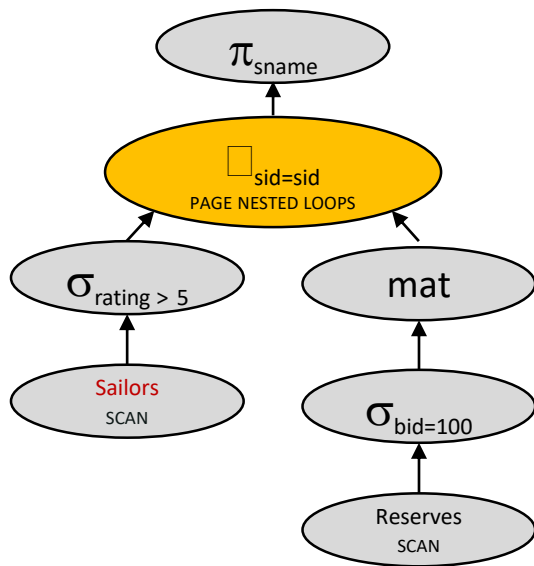


Query Plan 7 Cost Analysis Part 2

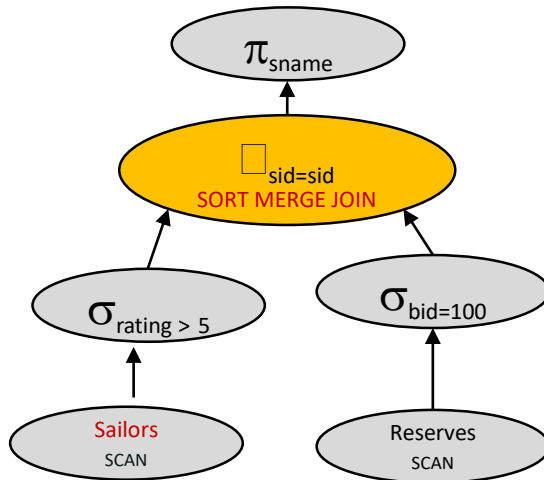
- With 5 buffers, cost of plan:
 - Scan Reserves (1000)
 - Scan Sailors (500)
 - Sort
 - 2 passes for reserves
pass 0 = 10 to write, pass 1 = 2*10 to read/write
 - 4 passes for sailors
pass 0 = 250 to write, pass 1,2,3 = 2*250 to read/write
 - Merge (10+250) = 260
- 1000 + 500 + sort reserves(10 + 2*10) + sort sailors
(250 + 3*2*250) + merge (10+250) = **3540**



Decision 5

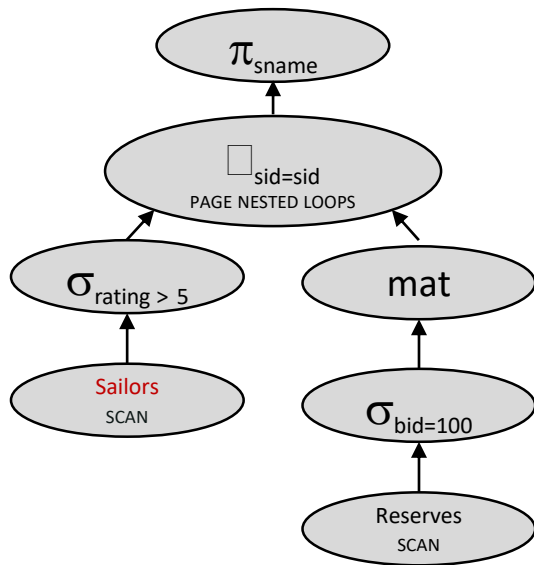


4010 IOs

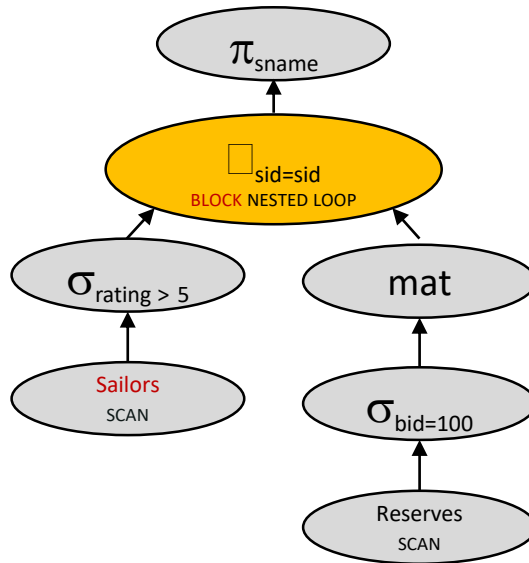


3540 IOs

Join Algorithm Again, Again



4010 IOs

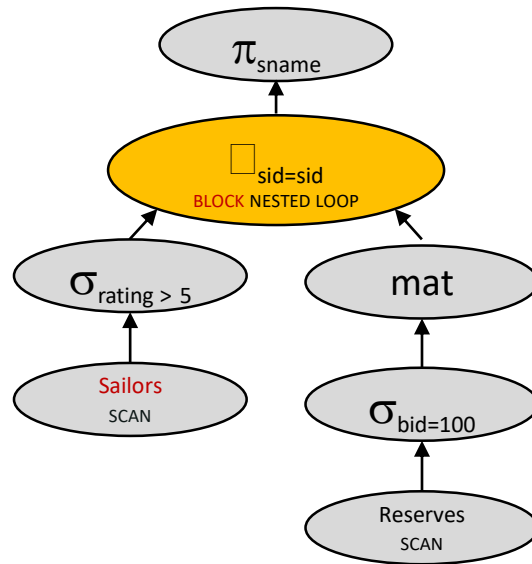


Cost???


Query 9 Cost Analysis

- With 5 buffers, cost of plan:
- Scan Sailors (500)
- Scan Reserves (1000)
- Write Temp T1 (10)
- For each blockful of high-rated sailors
 - Loop on T1 (??? * 10)
- Total:


$$500 + 1000 + 10 + (\text{ceil}(250/3) * 10) = 500 + 1000 + 10 + (84 * 10) =$$



Decision 7

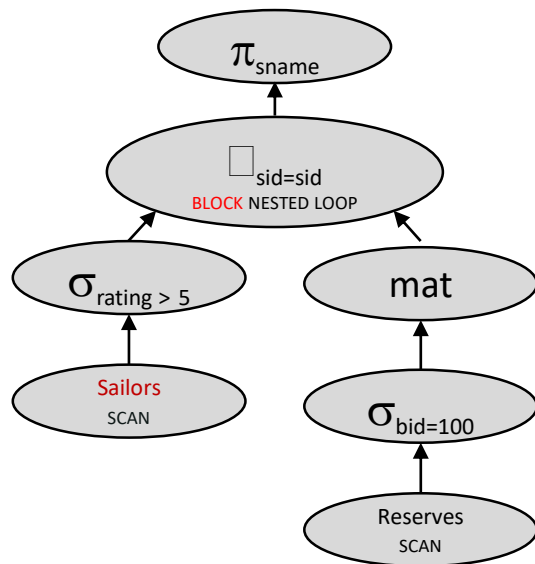


3540 IOs



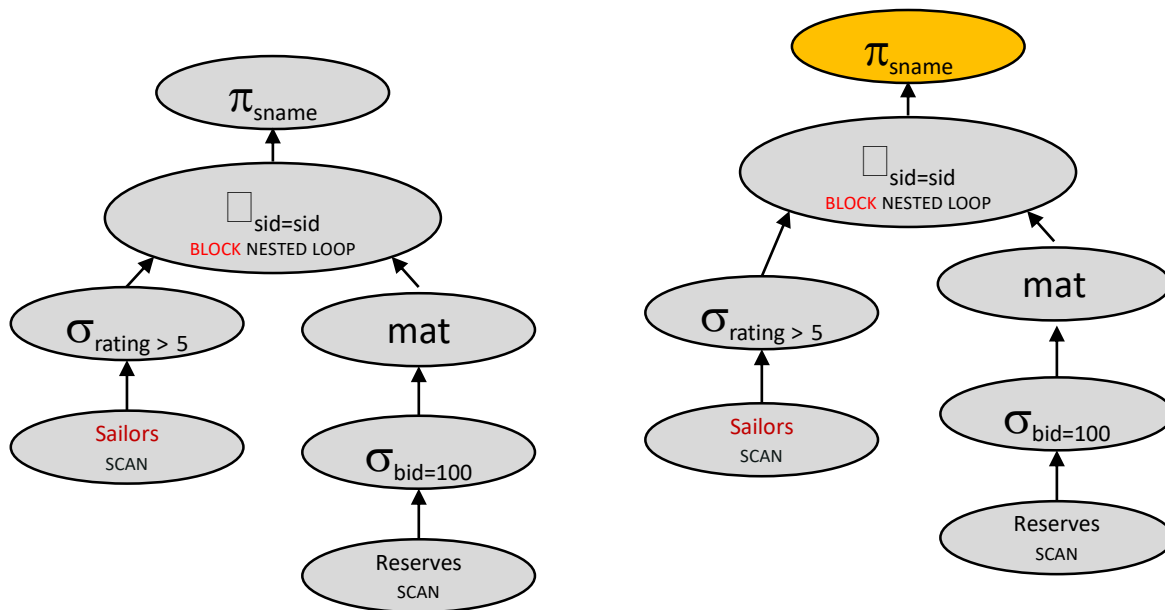
2350 IOs

Projection Cascade & Pushdown



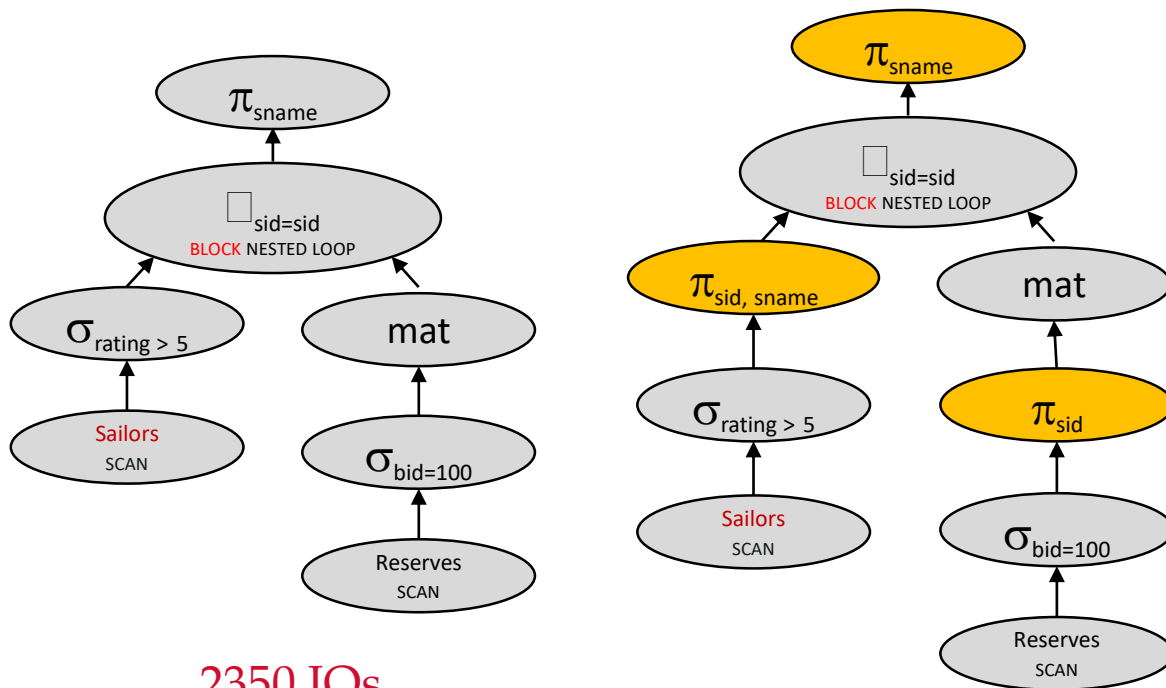
2350 IOs

Projection Cascade & Pushdown, cont



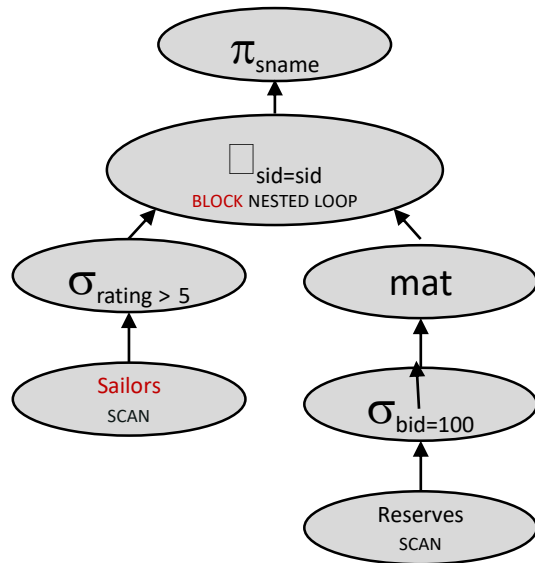
2350 IOs

Projection Cascade & Pushdown, cont

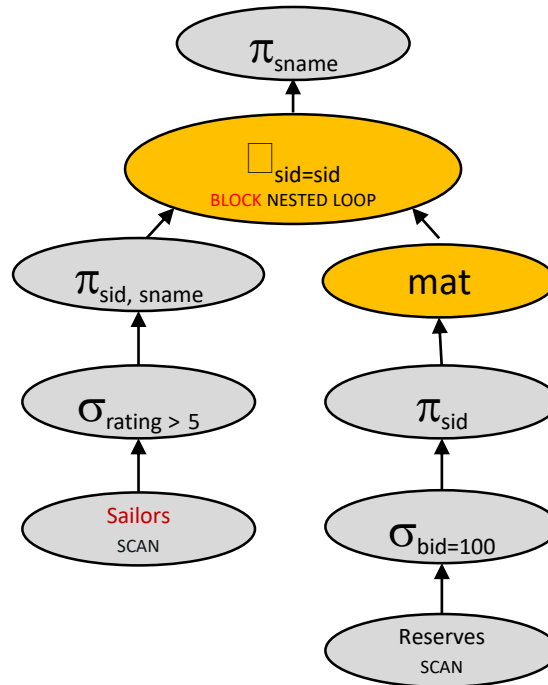


2350 IOs

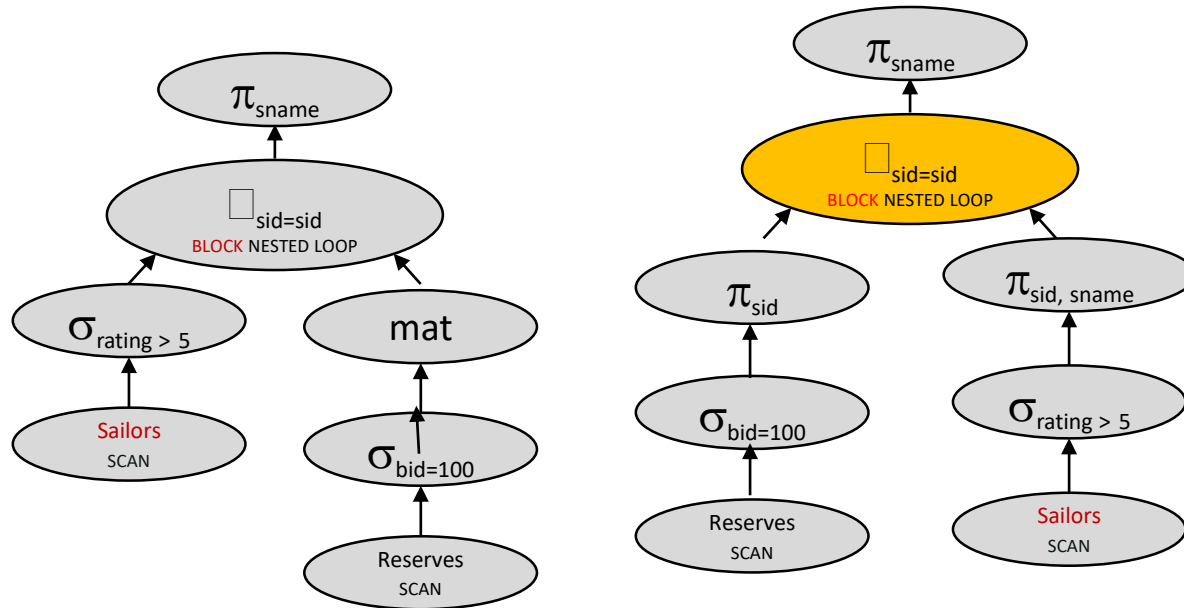
With Join Reordering, no Mat



2350 IOs



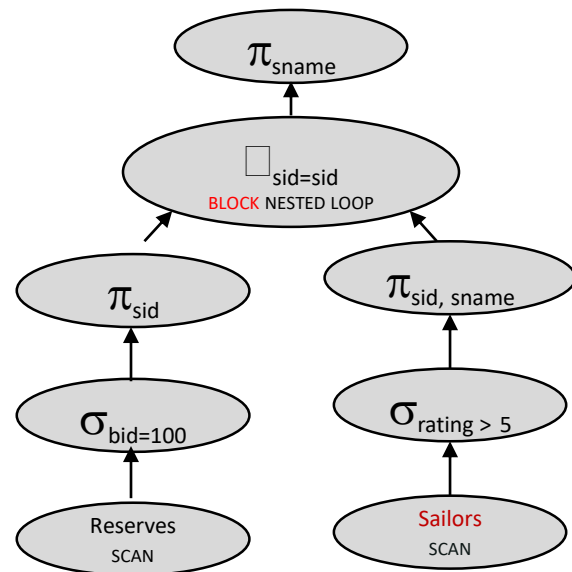
With Join Reordering, no Mat cont



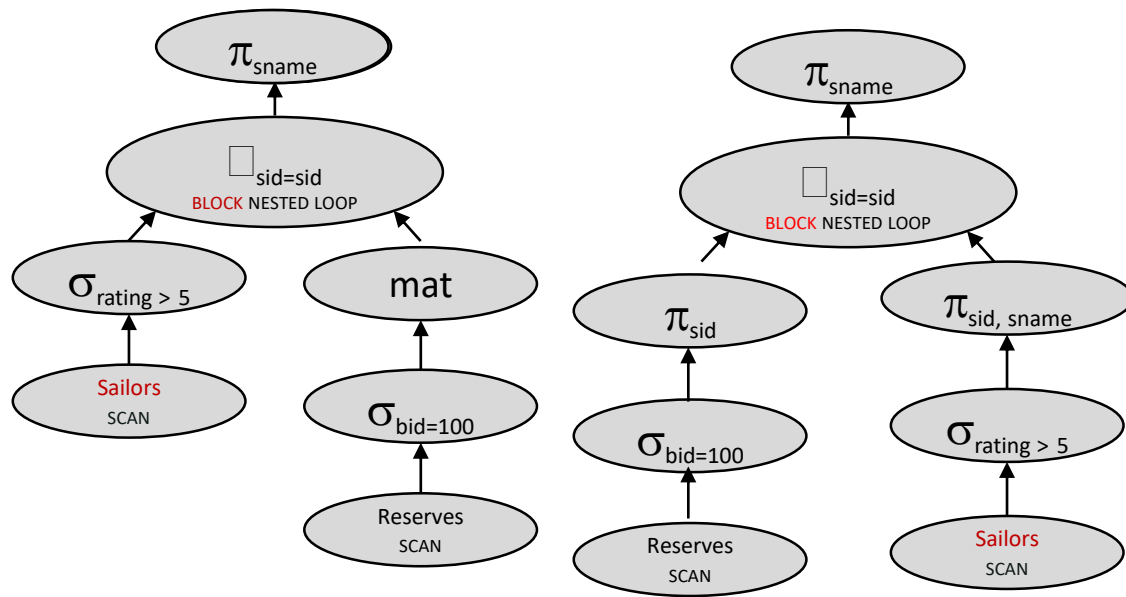
2350 IOs

Plan 11 Cost Analysis

- With 5 buffers, cost of plan:
- Scan Reserves (1000)
- For each blockful of sids that rented boat 100
- (recall Reserve tuple is 40 bytes, assume sid is 4 bytes)
- Loop on Sailors (??? * 500)
- Total: **1500**



With Join Reordering, no Mat, cont.



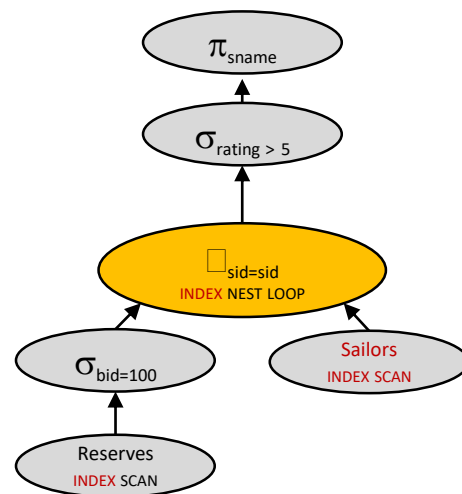


2350 IOs

1500 IOs

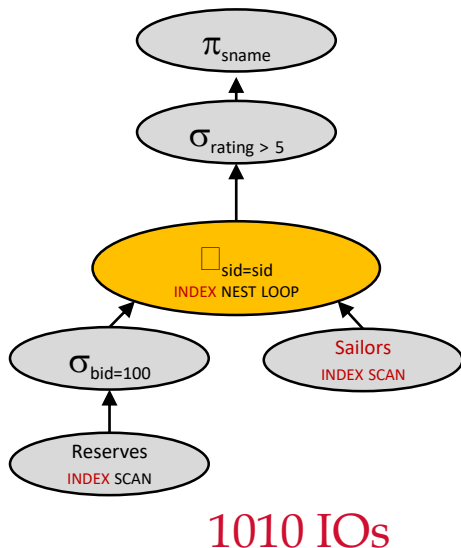
How About Indexes?

- Indexes:
 - Reserves.bid clustered
 - Sailors.sid unclustered
- Assume indexes fit in memory



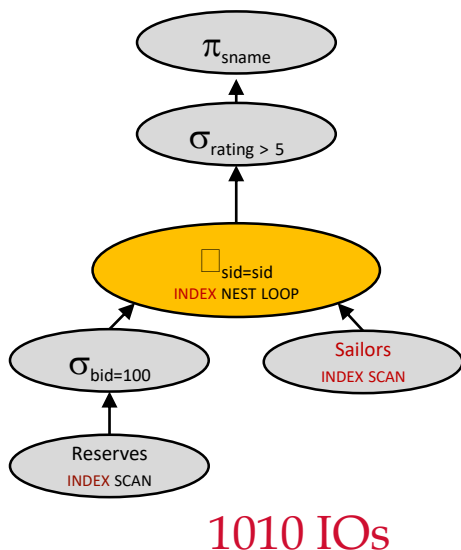
Index Cost Analysis

- **No projection pushdown to left** for π_{sname}
 - Projecting out unnecessary fields from outer of Index NL doesn't make an I/O difference.
- **No selection pushdown to right** for $\sigma_{rating > 5}$
 - Does not affect Sailors.sid index lookup
- With clustered index on bid of Reserves, we access how many pages of Reserves?:
 - $100,000/100 = 1000$ tuples on $1000/100 = 10$ pages.
- Join column sid is a **key** for Sailors.
 - At most one matching tuple, unclustered index on sid OK



Index Cost Analysis Part 2

- With clustered index on bid of Reserves, we access how many pages of Reserves?:
 - $100,000/100 = 1000$ tuples on $1000/100 = 10$ pages.
- for each Reserves tuple 1000
 get matching Sailors tuple (1 IO)
 (recall: 100 Reserves per page, 1000 pages)
- $10 + 1000*1$
- Cost: Selection of Reserves tuples (10 I/Os); then, for each, must get matching Sailors tuple (1000); total 1010 I/Os.



Summing up

- There are *lots* of plans
 - Even for a relatively simple query
- Engineers often think they can pick good ones
 - E.g. MapReduce API was based on that assumption
 - So was the COBOL API of 1970's!
- Not so clear that's true!
 - Manual query planning can be tedious, technical
 - Machines are better at enumerating options than people
 - Hence AI
 - We will see soon how optimizers make simplifying assumptions